

# Homework 2

Felipe Toledo  
ft8bn

March 16, 2021

## Clarifications

- I discussed the homework with Dane Williamson (dw3zn).

### Problem 1

Even though the softmax function is a nonlinear function, it can still be considered as a linear classifier. A linear classifier is one where a hyperplane is formed by taking a linear combination of the features, such that one side of the hyperplane predicts one class and the other side predicts the other. For multi-class classification we use the softmax function which calculates the probability of an input of being in a particular output class.

We can convert this multi-class classification problem into a binary classification problem if we separate one class from the rest of the classes, also known as "one versus all classification". In general, to solve a multi-class classification problem with  $n$  classes,  $n$  logistic regression models are trained each considering only a single class. Whenever a new data point has to be classified, the  $n$  models predict the probability of the point falling into its respective class. Finally, the point is assigned to the class with the highest probability as predicted by one of the  $n$  models.

In that case we can define two probabilities:  $P(y|x)$  and  $P(\neg y|x)$ . For binary classification, we know that those probabilities are:

$$P(y|x) = \frac{1}{1 + \exp(\boldsymbol{\omega}^T \mathbf{x} + b)} \quad (1)$$

$$P(\neg y|x) = 1 - P(y|x) = \frac{\exp(\boldsymbol{\omega}_1^T \mathbf{x} + b)}{1 + \exp(\boldsymbol{\omega}^T \mathbf{x} + b)} \quad (2)$$

The output would be of class "y" if the following is satisfied:

$$P(y|x) > P(\neg y|x)$$

which is equivalent to:

$$\frac{P(y|x)}{P(\neg y|x)} > 1$$

If we apply the logarithm to both sides, we will get:

$$\begin{aligned} \log\left(\frac{P(y|x)}{P(\neg y|x)}\right) &> \log(1) \\ \log\left(\frac{P(y|x)}{P(\neg y|x)}\right) &> 0 \end{aligned} \quad (3)$$

If we replace (1) and (2) into (3) we will get:

$$\begin{aligned} \log(\exp(\boldsymbol{\omega}_1^T \mathbf{x} + b)) - \log(1 + \exp(\boldsymbol{\omega}^T \mathbf{x} + b)) - \log(1) + \log(1 + \exp(\boldsymbol{\omega}^T \mathbf{x} + b)) &> 0 \\ \log(\exp(\boldsymbol{\omega}_1^T \mathbf{x} + b)) - \emptyset &> 0 \\ \boldsymbol{\omega}_1^T \mathbf{x} + b &> 0 \end{aligned} \quad (4)$$

From (4) we see that the decision boundary is given by the plane  $\omega_1^T \mathbf{x} + b$ , consequently we can conclude that the softmax function is still a linear classifier.

## Problem 2

Lets start from the definition of the softmax function:

$$P(y|x) = \frac{\exp(\omega_y^T \mathbf{x})}{\sum_{y' \in Y} \exp(\omega_{y'}^T \mathbf{x})}$$

Now lets find out what is the  $P(y = 1|x)$

$$\begin{aligned} P(y = 1|x) &= \frac{\exp(\omega_1^T \mathbf{x})}{\exp(\omega_0^T \mathbf{x}) + \exp(\omega_1^T \mathbf{x})} \\ &= \frac{\exp(\omega_1^T \mathbf{x})}{\exp(\omega_0^T \mathbf{x}) + \exp(\omega_1^T \mathbf{x})} \cdot \frac{\exp(-\omega_1^T \mathbf{x})}{\exp(-\omega_1^T \mathbf{x})} \\ &= \frac{\exp(\omega_1^T \mathbf{x} - \omega_1^T \mathbf{x})}{\exp(\omega_0^T \mathbf{x} - \omega_1^T \mathbf{x}) + \exp(\omega_1^T \mathbf{x} - \omega_1^T \mathbf{x})} \\ &= \frac{\exp(0)}{\exp((\omega_0^T - \omega_1^T) \mathbf{x}) + \exp(0)} \\ &= \frac{1}{\exp((\omega_0^T - \omega_1^T) \mathbf{x}) + 1} \end{aligned} \tag{1}$$

Softmax function is used for multi-class classification problems, in the above equation we end up having  $\omega_0^T$  and  $\omega_1^T$  that represent the classification weight associated with the two possible labels. To simplify the equation, we can express the subtraction of the weights in the following way:

$$\begin{aligned} \omega^T &= \omega_0^T - \omega_1^T \\ -\omega^T &= \omega_1^T - \omega_0^T \end{aligned} \tag{2}$$

If we replace (2) in (1), we get:

$$P(y = 1|x) = \frac{1}{\exp(-\omega^T \mathbf{x}) + 1}$$

## Problem 3

Lets start by the definition of cross-entropy loss function:

$$H(Q(Y|x), P(Y|x)) = - \sum_{y' \in Y} Q(Y = y'|x) \log P(Y = y'|x) \tag{3}$$

With the empirical distribution:

$$Q(Y = y'|x^{(i)}) = \begin{cases} 1 & y' = y \\ 0 & y' \neq y \end{cases}$$

Based on the empirical distribution, we can see in (1) that we are only going to sum the terms in which the predicted label  $y'$  is equal to the ground truth  $y$ . Moreover, if we use the collection of training examples denoted by  $\{(x^{(i)}, y^{(i)})\}_{i=1}^m$  instead of the predicted labels  $y' \in Y$ . We can replace  $Q(Y = y'|x^{(i)})$  in (1) by the ground truth probability of  $y^{(i)}$  which will always be 1 according to the empirical distribution, and  $P(Y = y'|x)$  by the probability of the predicted classes  $P(Y = y^{(i)}|x^{(i)})$ . By doing that, we will need to modify the bounds of the summation to iterate over all the training data. Obtaining:

$$\sum_{i=1}^m H(Q(Y|x^{(i)}), P(Y|x^{(i)})) = - \sum_{i=1}^m y^{(i)} \log P(Y = y^{(i)}|x^{(i)})$$

$$\sum_{i=1}^m H(Q(Y|x^{(i)}), P(Y|x^{(i)})) = L(\theta)$$

## Problem 4

The code for this problem is in a python notebook called "ft8bn-q4.ipynb". The results of the hypertunning of the parameters was saved in a file called "results.pkl", this file is loaded in the same notebook. As indicated in point (e) the test headlines predictions were saved in a file called "news-tst.pred".

## Problem 5

The code for this problem is in a python notebook called "ft8bn-cnn.ipynb". The validation accuracy that I got from the cnn model is 0.6352.